

Deep Learning Project - Hieroglyphs

MBA Tech CE - 3rd Year

Project by:

Swapnil Singh (N041)

Vaishnavi Singh (N048)

Tarun Tanmay (N049)

Vidhi Vaziani (N052)

Introduction

- This project is a small part of a larger goal
- The aim of the project is to classify egyptian symbols or Hieroglyphs into the classes assigned by the Archeologist
- Other sections of the project including
 - Image processing to capture, filter, and segment the image to extract the individual symbols
 - Followed by this section where we classify the image to the annotations assigned to it
 - Then comes the concepts of NLP where we aim to translate the annotation to normal english for people to read
- Given below is an image from the Pyramid of Giza





Loading the dataset

1. Accessing the drive for images
2. Loading the images
3. Making batches

Vidhi Vazirani - N052

In []:

```
from google.colab import drive
drive.mount("/content/drive", force_remount=True)
```

Mounted at /content/drive

In []:

```
import numpy as np
from multiprocessing.pool import Pool
from keras.preprocessing import image
from keras.applications.inception_v3 import preprocess_input
```

In []:

```
def loadImage(path):
    img = image.load_img(path, target_size=(299, 299))
    x = image.img_to_array(img)
    x = np.expand_dims(x, axis=0)
    x = preprocess_input(x)
    return x
```

```
In [ ]:
```

```
def loadBatch(img_paths):  
    with Pool(processes=8) as pool:  
        imgs = pool.map(loadImage, img_paths)  
        return np.vstack(imgs)
```

```
In [ ]:
```

```
def batchGenerator(img_paths, labels, batch_size):  
    for i in range(0, len(img_paths), batch_size):  
        batch_paths = img_paths[i:(i + batch_size)]  
        batch_labels = labels[i:(i + batch_size)]  
        batch_images = loadBatch(batch_paths)  
        yield batch_images, batch_labels
```

Feature Extraction

Feature extraction is done using the Inception network.

The inception network works in 2 steps: Feature extraction and classification. But here, we just need the features. So we pop out the last layer of the network.

```
In [ ]:
```

```
from keras.applications.inception_v3 import InceptionV3  
from sklearn.preprocessing import normalize
```

```
In [ ]:
```

```
class FeatureExtractor:  
    def __init__(self):  
        print("loading DeepNet (Inception-V3) ...")  
        self.model = InceptionV3(weights='imagenet')  
  
        # Initialise the model to output the second to last layer, which contains the de  
        eplearning featuers  
        self.model.layers.pop() # Get rid of the classification layer  
        self.model.outputs = [self.model.layers[-1].output]  
        self.model.layers[-1].outbound_node = []  
  
    def get_features(self, batch):  
        features = self.model.predict(batch)  
        features = features.reshape(-1, features.shape[-1])  
        return normalize(features, axis=1, norm='l2')
```

```
In [ ]:
```

```
import numpy as np  
from os import listdir, path  
from os.path import isdir, isfile, join, exists, dirname  
from sklearn.model_selection import train_test_split, GridSearchCV  
from sklearn import linear_model  
from sklearn.externals import joblib  
from urllib.request import urlopen  
from zipfile import ZipFile  
from io import BytesIO  
import os  
import cv2
```

```
/usr/local/lib/python3.7/dist-packages/sklearn/externals/joblib/__init__.py:15: FutureWar  
ning: sklearn.externals.joblib is deprecated in 0.21 and will be removed in 0.23. Please  
import this functionality directly from joblib, which can be installed with: pip install  
joblib. If this warning is raised when loading pickled models, you may need to re-seriali  
ze those models with scikit-learn 0.21+.  
warnings.warn(msg, category=FutureWarning)
```

In []:

```
dataPath      = '/content/drive/MyDrive/Dataset'
stelePath     = join(dataPath, "Manual/Preprocessed")
examplePath   = join(dataPath, "Examples")
featurePath   = "features.npy"
labelsPath    = "labels.npy"
svmPath       = "svm.pkl"
image_paths   = []
labels        = []
batch_size    = 2000
```

In []:

```
print("indexing images...")
Steles = [ join(stelePath,f) for f in listdir(stelePath) if isdir(join(stelePath,f)) ]
for stele in Steles:
    imagePaths = [ join(stele,f) for f in listdir(stele) if isfile(join(stele,f)) ]
    for path in imagePaths:
        image_paths.append(path)
        labels.append(path[(path.rfind("_") + 1): path.rfind(".")] )

featureExtractor = FeatureExtractor()
features = []
print("computing features...")
for idx, (batch_images, _) in enumerate(batchGenerator(image_paths, labels, batch_size)):
    :
    print("{} / {}".format((idx+1) * batch_size, len(labels)))
    features_ = featureExtractor.get_features(batch_images)
    features.append(features_)
features = np.vstack(features)

labels = np.asarray(labels)
print("saving features...")
np.save(featurePath, features)
np.save(labelsPath, labels)
```

```
indexing images...
loading DeepNet (Inception-V3) ...
Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/inception_v3/inception_v3_weights_tf_dim_ordering_tf_kernels.h5
96116736/96112376 [=====] - 1s 0us/step
computing features...
2000/4210
4000/4210
6000/4210
saving features...
```

In []:

```
features
```

Out[]:

```
array([[1.5229225e-04, 5.5373140e-04, 3.0603906e-04, ..., 2.1803525e-04,
        1.3289934e-04, 6.4251979e-04],
       [6.4444571e-04, 4.5223985e-04, 5.2867125e-04, ..., 5.3320249e-04,
        3.7295202e-04, 2.8969211e-04],
       [2.0304059e-04, 2.7304085e-04, 3.5118556e-03, ..., 9.1406232e-04,
        4.3884842e-04, 4.9880863e-04],
       ...,
       [5.6005852e-06, 2.2427563e-05, 8.1763294e-04, ..., 1.7655144e-05,
        3.6304511e-05, 1.3711563e-05],
       [1.7347759e-04, 2.4198461e-04, 5.2247121e-04, ..., 1.5125928e-03,
        1.8769264e-04, 2.7925041e-04],
       [8.8263077e-05, 1.8308988e-04, 7.6894701e-04, ..., 4.7692665e-05,
        1.2536222e-04, 1.5528710e-04]], dtype=float32)
```

Filtering the dataset

Removing images that do not have a class associated with it.

Splitting the dataset into training and testing samples (80%-20%)

In []:

```
tobeDeleted = np.nonzero(labels == "UNKNOWN")
features = np.delete(features,tobeDeleted, 0)
labels = np.delete(labels,tobeDeleted, 0)
numImages = len(labels)
trainSet, testSet, trainLabels, testLabels = train_test_split(features, labels, test_size=0.20, random_state=42)
```

Training basic Classifiers

Training some machine learning classifiers like Logistic Regression and XGBoost.

Training a deep learning classifier - MLPClassifier.

The above classifiers yield low accuracy scores.

In []:

```
#Logistic regression is a classification algorithm, used when the value of the target variable is categorical in nature.
# most commonly used when the output variable is categorical in nature
print("training SVM...")
if 0: # optional; either train 1 classifier fast, or search trough the parameter space by training multiple classifiers to squeeze out that extra 2%
    clf = linear_model.LogisticRegression(C=10000) #This class implements regularized logistic regression
else:
    svr = linear_model.LogisticRegression(max_iter=10000)
    parameters = {'C':[0.001, 0.01, 0.1, 1, 10, 100, 1000, 10000]}
    clf = GridSearchCV(svr, parameters, n_jobs=8) #Exhaustive search over specified parameter values for an estimator.
#Important members are fit, predict.
#GridSearchCV implements a "fit" and
a "score" method
clf.fit(trainSet, trainLabels) # fitting the training set to the above classifier

print(clf)
print("finished training! saving...")
joblib.dump(clf, 'clf.pkl', compress=1) # saving the model

prediction = clf.predict(testSet) #predict() function enables us to predict the labels of the data values on the basis of the trained model.
accuracy = np.sum(testLabels == prediction) / float(len(prediction)) # calculating the accuracy reflected by the above classifier manually

# for idx, pred in enumerate(prediction):
#     print("%-5s --> %s" % (testLabels[idx], pred))
print("accuracy = {}".format(accuracy*100))
```

training SVM...

```
/usr/local/lib/python3.6/dist-packages/sklearn/model_selection/_split.py:667: UserWarning
: The least populated class in y has only 1 members, which is less than n_splits=5.
% (min_groups, self.n_splits)), UserWarning)
```

```
GridSearchCV(cv=None, error_score=nan,
             estimator=LogisticRegression(C=1.0, class_weight=None, dual=False,
```



```

        fit_intercept=True,
        intercept_scaling=1, l1_ratio=None,
        max_iter=10000, multi_class='auto',
        n_jobs=None, penalty='l2',
        random_state=None, solver='lbfgs',
        tol=0.0001, verbose=0,
        warm_start=False),

```

```

        iid='deprecated', n_jobs=8,
        param_grid={'C': [0.001, 0.01, 0.1, 1, 10, 100, 1000, 10000]},
        pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
        scoring=None, verbose=0)

```

finished training! saving...
accuracy = 64.56009913258984%

In []:

```
from xgboost import XGBClassifier
```

In []:

```

#XGBoost is a powerful machine learning algorithm especially where speed and accuracy are
concerned
xgb = XGBClassifier()
xgb.fit(trainSet, trainLabels) # fitting the training set to the above classifier

joblib.dump(xgb, svmPath, compress=1) # saving the model

prediction = xgb.predict(testSet) #predict() function enables us to predict the labels of
the data values on the basis of the trained model.
accuracy = np.sum(testLabels == prediction) / float(len(prediction)) # calculating the
accuracy reflected by the above classifier manually

# for idx, pred in enumerate(prediction):
#     print("%-5s --> %s" % (testLabels[idx], pred))
print("accuracy = {}".format(accuracy*100))

```

accuracy = 68.02973977695167%

In []:

```
from sklearn.neural_network import MLPClassifier
```

In []:

```

#Multi-layer Perceptron classifier.

# multilayer perceptron (MLP) is a class of feedforward artificial neural network (ANN)
mlp = MLPClassifier(max_iter=1000, hidden_layer_sizes=(1000, 1000, 500, 171) , alpha=0.0
001, solver='adam') #alpha= regularization term ,adam' refers to a stochastic gradient-b
ased optimizer
#The default solver 'adam' works pretty well on relatively large datasets
#(with thousands of training samples or more) in terms of both training time and validati
on score.
# For small datasets, however, 'lbfgs' can converge faster and perform better.
mlp.fit(trainSet, trainLabels)

joblib.dump(mlp, "mlp.pkl", compress=1) # saving

prediction = mlp.predict(testSet) #predict() function enables us to predict the labels of
the data values on the basis of the trained model.
accuracy = np.sum(testLabels == prediction) / float(len(prediction)) # calculating accur
acy manually

# for idx, pred in enumerate(prediction):
#     print("%-5s --> %s" % (testLabels[idx], pred))
print("accuracy = {}".format(accuracy*100))

```

accuracy = 64.80793060718712%

In []:

```
import os
```

Testing XGBoost

Since XGBoost gave the highest accuracy, as seen above, we can test using some samples and check the accuracy.

In []:

```
inputPath = examplePath
if isdir(inputPath):
    imagePaths = [join(inputPath, f) for f in listdir(inputPath) if f.endswith(('.png',
    '.jpg'))]
else:
    imagePaths = [inputPath,]

print("loading images...")
Images = loadBatch(imagePaths)
print("loading SVM model...")
clf = joblib.load(svmPath);

print("Extracting features, this may take a while for large collections of images...") #
should probably use batches for this as well
extractor = FeatureExtractor() #This class allows you to extract features of an image via
a pre-trained model and re-train that model with new data.

features = extractor.get_features(Images) #Extract set of features from a time series(m
atrix of dataframe).

classes = xgb.classes_
print("Predicting the Hieroglyph type...")
prob = np.array(xgb.predict_proba(features))
top5_i = np.argsort(-prob)[: ,0]
top5_s = np.array([prob[row, top5_i[row]] for row, top5_i_row in enumerate(top5_i)])
top5_n = classes[top5_i]

print("{:<25} ::: {}".format("image name", "top 5 best matching hieroglyphs"))
for idx, path in enumerate(imagePaths):
    print("{:<25} --> {}".format(os.path.basename(path), top5_n[idx]))
```

```
loading images...
loading SVM model...
Extracting features, this may take a while for large collections of images...
loading DeepNet (Inception-V3) ...
Predicting the Hieroglyph type...
image name          ::: top 5 best matching hieroglyphs
200000_S29.png      --> S29
200001_V13.png      --> V13
200002_V13.png      --> V13
200003_G43.png      --> G43
200004_D21.png      --> D21
200005_O50.png      --> O50
200006_X1.png       --> X1
200007_M23.png      --> M23
200008_G43.png      --> G43
200009_S29.png      --> S29
200010_V13.png      --> V13
200011_M23.png      --> M23
200012_G43.png      --> G43
200013_D21.png      --> D21
200014_O50.png      --> O50
200015_V13.png      --> V13
200016_G43.png      --> G43
```

Preparing data for CNN

Swapnil Singh - N041

In []:

```
# importing the required libraries for reading the images
from keras.preprocessing.image import ImageDataGenerator
import cv2
```

Augmenting the data

In []:

```
# defining a function for augmenting and storing the images.
def augment_data(file_dir, n_generated_samples, save_to_dir, taking):

    data_gen = ImageDataGenerator(rotation_range=10,
                                   width_shift_range=0.1,
                                   height_shift_range=0.1,
                                   shear_range=0.1,
                                   brightness_range=(0.3, 1.0),
                                   horizontal_flip=True,
                                   vertical_flip=True,
                                   fill_mode='nearest'
                                   )

    # defining the type of augmentation we want.
    # we defined a rotation range of 10, width shift of 0.1, height shift of 0.1, shear range
    # of 0.1, brightness range of 0.3-1.0
    # horizontal flip and vertical flips were set to true.

    for filename in taking:
        # load the image
        image = cv2.imread(file_dir + '/' + filename)
        # reshape the image
        image = image.reshape((1,)+image.shape)
        # prefix of the names for the generated sampels.
        save_prefix = 'aug_' + filename[:filename.rfind('.')]
        print(save_prefix)
        # generate 'n_generated_samples' sample images
        i=0
        for batch in data_gen.flow(x=image, batch_size=1, save_to_dir=save_to_dir,
                                   save_prefix=save_prefix, save_format='png
        '):
            i += 1
            if i > n_generated_samples:
                break

    # in the above code we are reading the images, resizing them, and then augmenting and sav
    # ing the images with the prefic aug_
```

In []:

```
# DON'T RUN THIS CELL AGAIN.
augmented_data_path = '/content/drive/MyDrive/GlyphDataset/Augmented_hieroglyphs'
# declaring the augmentation path

# augment data for the examples with label equal to 'yes' representing tumorous examples
augment_data(file_dir='/content/drive/MyDrive/GlyphDataset/hieroglyphs', n_generated_sam
ples=10, save_to_dir=augmented_data_path, taking=taking)
# calling the function for saving the augmented images in the given path, here we have ma
de 10 images for each image passed
```


In []:

```
# declaring paths
image_paths1      = []
labels1           = []
features1         = []
batch_size        = 2000
```

Manually Preprocessed Data

In []:

```
# storing the path for images and storing lables in the declared arrays
print("indexing images...")
Steles = [ join(stelePath,f) for f in listdir(stelePath) if isdir(join(stelePath,f)) ]
for stele in Steles:
    imagePaths1 = [ join(stele,f) for f in listdir(stele) if isfile(join(stele,f)) ]
    for path in imagePaths1:
        image_paths1.append(path)
        #print(path)
        labels1.append(path[(path.rfind("_") + 1): path.rfind(".")]])
for filename in os.listdir("/content/drive/MyDrive/Dataset/Augmented"):
    image_paths1.append(join('/content/drive/MyDrive/Dataset/Augmented', filename))
    labels1.append(path[(path.rfind("_") + 1): path.rfind(".")]])
print(len(labels1))
```

indexing images...
63680

In []:

```
# storing the path for images and storing lables in the declared arrays for the second te
st dataset
print("indexing images...")
image_paths2 = []
labels2 = []
for filename in taking:
    image_paths2.append(join('/content/drive/MyDrive/GlyphDataset/hieroglyphs', filename
))
    labels2.append(filename[:filename.rfind(".")])
for filename in os.listdir('/content/drive/MyDrive/GlyphDataset/Augmented_hieroglyphs'):
    image_paths2.append(join('/content/drive/MyDrive/GlyphDataset/Augmented_hieroglyphs'
, filename))
    labels2.append(filename[4 : (filename.rfind("_")-2)])
print(len(labels2))
```

indexing images...
1967

In []:

```
taking = []
```

In []:

```
# adding non augmented images to the dataset
for filename in os.listdir('/content/drive/MyDrive/GlyphDataset/hieroglyphs'):
    if filename[:filename.rfind('.')] in labels1:
        taking.append(filename)
```

In []:

```
len(np.unique(labels1))
```

Out[]:

171

In []:

```
# removing images whose class is not known
tobeDeleted = np.nonzero(labels1 == "UNKNOWN") # Remove the Unknown class from the database
image_paths1 = np.delete(image_paths1,tobeDeleted, 0)
labels1 = np.delete(labels1,tobeDeleted, 0)
numImages = len(labels1)
```

In []:

```
# label encoding and then transforming the labels.
label_encoder = LabelEncoder()
label_encoder.fit(labels1)
labels2_encoded = label_encoder.transform(labels2)
```

In []:

```
# checking the encoded labels
labels2_encoded
```

Out[]:

```
array([150, 143, 145, ..., 81, 81, 81])
```

In []:

```
# checking the classes in encoded dataset
np.unique(labels2_encoded)
```

Out[]:

```
array([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11, 12,
        13, 14, 15, 16, 17, 18, 19, 20, 21, 22, 23, 24, 25,
        26, 27, 28, 29, 30, 31, 32, 33, 34, 35, 36, 37, 38,
        39, 40, 41, 42, 43, 44, 45, 46, 47, 48, 49, 50, 51,
        52, 53, 54, 55, 56, 57, 58, 59, 60, 61, 62, 63, 64,
        65, 66, 67, 68, 69, 70, 71, 72, 73, 74, 76, 77, 78,
        79, 80, 81, 82, 84, 85, 86, 87, 88, 89, 90, 91, 92,
        93, 94, 95, 96, 97, 99, 100, 101, 102, 104, 105, 106, 107,
       108, 110, 111, 112, 113, 114, 115, 116, 117, 118, 120, 121, 122,
       123, 124, 125, 126, 127, 128, 129, 131, 132, 133, 134, 135, 136,
       137, 138, 139, 140, 141, 142, 143, 144, 145, 146, 147, 148, 149,
       150, 151, 152, 153, 154, 155, 156, 157, 158, 159, 160, 161, 162,
       163, 164, 165, 166, 167, 168, 169, 170])
```

In []:

```
# manually one hot encoding images for the second dataset
labels2_one_hot_encoded = []
for i in labels2_encoded:
    a = np.zeros(171,)
    a[i] = 1
    labels2_one_hot_encoded.append(a)
```

In []:

```
# verifying the count of the classes
target = label_encoder.classes_
len(target)
```

Out[]:

```
171
```

In []:

```
# one hot encoding the labels of the training dataset
```

```
labels1 = to_categorical(labels1)
```

In []:

```
#iterating in training set of data and reading and resizing the images
rawImages2 = []
for i in image_paths2:
    img = cv2.imread(i, cv2.IMREAD_GRAYSCALE)
    #stores the raw pixel values of this image after resizing
    pixels = getPixels(img, (32,32))
    #stores the raw pixel values of images
    rawImages2.append(pixels)
```

In []:

```
rawImages2 = np.asarray(rawImages2) # converting to numpy array
labels2_one_hot_encoded = np.asarray(labels2_one_hot_encoded) # converting to numpy array
```

In []:

```
labels2_one_hot_encoded.shape # checking the shape of the encoded labels
```

Out[]:

```
(1967, 171)
```

In []:

```
np.save('/content/drive/MyDrive/Dataset/rawimages_ejypt.npy', rawImages1) # saving the dataset read to save time in the future.
# next time use np.load('/content/drive/MyDrive/Dataset/rawimages_ejypt.npy')
```

Convolution Neural Network (CNN)

Tarun Tanmay - N049

In []:

```
from keras.utils import to_categorical
```

Preparing Training and testing data

In []:

```
X_train1,X_test1,y_train1,y_test1 = train_test_split(rawImages1,labels1,test_size = 0.2,
random_state=42)
```

In []:

```
X_train1 = X_train1.reshape(-1, 32, 32, 1)
X_test1 = X_test1.reshape(-1, 32, 32, 1)
```

In []:

```
rawImages2 = rawImages2.reshape(-1, 32, 32, 1)
rawImages2 = rawImages2 / 255.0
```

In []:

```
X_train1 = X_train1 / 255.0
X_test1 = X_test1 / 255.0
```

In []:

```
from sklearn.preprocessing import OneHotEncoder, LabelEncoder
```

In []:

```
from tensorflow.keras import layers, models, utils, datasets
from keras.layers import Dense, Flatten, Conv2D, MaxPooling2D, Dropout, BatchNormalizati
on
from keras.layers import LeakyReLU
```

Custom CNN

In []:

```
#creating A Convolutional Neural Network
model1=models.Sequential() # sequential neural netwo
rk
model1.add(Conv2D(32, (5, 5), input_shape = (32, 32, 1))) # 5*5 size filter, input l
ayer has 32 neurons
model1.add(LeakyReLU(alpha=0.1)) # alpha means learning rat
e
model1.add(MaxPooling2D(pool_size = (2, 2))) # max-pooling
model1.add(Conv2D(128, (5, 5)))
model1.add(LeakyReLU(alpha=0.1))
model1.add(Conv2D(64, (5, 5)))
model1.add(LeakyReLU(alpha=0.1))
model1.add(Conv2D(32, (5, 5)))
model1.add(LeakyReLU(alpha=0.1))
model1.add(MaxPooling2D(pool_size = (2, 2)))
model1.add(Flatten()) # flattening o/p of cnn
model1.add(Dense(1000)) # Dense layer (Fully conn
ected layer)
model1.add(LeakyReLU(alpha=0.1))
model1.add(Dropout(0.5)) # Dropout to reducde overf
itting
model1.add(Dense(500))
model1.add(LeakyReLU(alpha=0.1))
model1.add(Dropout(0.5))
model1.add(Dense(250))
model1.add(LeakyReLU(alpha=0.1))
model1.add(Dense(171, activation = 'softmax')) # softmax activation funct
ion used for measuring predictions
#output layer has 171 neurons, which represnet the 171 classes of the dataset
model1.summary() # summarising what is in
the designed model
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 28, 28, 32)	832
leaky_re_lu (LeakyReLU)	(None, 28, 28, 32)	0
max_pooling2d (MaxPooling2D)	(None, 14, 14, 32)	0
conv2d_1 (Conv2D)	(None, 10, 10, 128)	102528
leaky_re_lu_1 (LeakyReLU)	(None, 10, 10, 128)	0
conv2d_2 (Conv2D)	(None, 6, 6, 64)	204864
leaky_re_lu_2 (LeakyReLU)	(None, 6, 6, 64)	0
conv2d_3 (Conv2D)	(None, 2, 2, 32)	51232

leaky_re_lu_3 (LeakyReLU)	(None, 2, 2, 32)	0
max_pooling2d_1 (MaxPooling2D)	(None, 1, 1, 32)	0
flatten (Flatten)	(None, 32)	0
dense (Dense)	(None, 1000)	33000
leaky_re_lu_4 (LeakyReLU)	(None, 1000)	0
dropout (Dropout)	(None, 1000)	0
dense_1 (Dense)	(None, 500)	500500
leaky_re_lu_5 (LeakyReLU)	(None, 500)	0
dropout_1 (Dropout)	(None, 500)	0
dense_2 (Dense)	(None, 250)	125250
leaky_re_lu_6 (LeakyReLU)	(None, 250)	0
dense_3 (Dense)	(None, 171)	42921
=====		
Total params: 1,061,127		
Trainable params: 1,061,127		
Non-trainable params: 0		

In []:

```
#we use the adam optimizer to handle sparse gradient for noisy problems, it combines the
benefits of rmsprop and adagrad algorithms
model1.compile(optimizer='adam',
               loss='categorical_crossentropy',
               metrics=['accuracy'])
               #we use categorical_entropy
               #If we use this loss, we will train a CNN to output a probability over the
               C classes for each image.
               #It is used for multi-class classification

history1 = model1.fit(X_train1, y_train1, epochs= 50, #training the model for 50 epochs
                    validation_split = 0.1)
# using the default batch size of 32
```

```
Epoch 1/50
1429/1429 [=====] - 14s 5ms/step - loss: 0.5253 - accuracy: 0.93
25 - val_loss: 0.3001 - val_accuracy: 0.9404
Epoch 2/50
1429/1429 [=====] - 7s 5ms/step - loss: 0.2916 - accuracy: 0.938
9 - val_loss: 0.2266 - val_accuracy: 0.9531
Epoch 3/50
1429/1429 [=====] - 7s 5ms/step - loss: 0.2191 - accuracy: 0.950
2 - val_loss: 0.2127 - val_accuracy: 0.9547
Epoch 4/50
1429/1429 [=====] - 7s 5ms/step - loss: 0.1977 - accuracy: 0.954
4 - val_loss: 0.1667 - val_accuracy: 0.9646
Epoch 5/50
1429/1429 [=====] - 7s 5ms/step - loss: 0.1604 - accuracy: 0.961
3 - val_loss: 0.1433 - val_accuracy: 0.9709
Epoch 6/50
1429/1429 [=====] - 7s 5ms/step - loss: 0.1458 - accuracy: 0.965
6 - val_loss: 0.1757 - val_accuracy: 0.9652
Epoch 7/50
1429/1429 [=====] - 7s 5ms/step - loss: 0.1448 - accuracy: 0.966
8 - val_loss: 0.1450 - val_accuracy: 0.9736
Epoch 8/50
1429/1429 [=====] - 7s 5ms/step - loss: 0.1315 - accuracy: 0.967
```

8 - val_loss: 0.1184 - val_accuracy: 0.9774
Epoch 9/50
1429/1429 [=====] - 7s 5ms/step - loss: 0.1394 - accuracy: 0.966
4 - val_loss: 0.1246 - val_accuracy: 0.9736
Epoch 10/50
1429/1429 [=====] - 7s 5ms/step - loss: 0.1221 - accuracy: 0.969
4 - val_loss: 0.1133 - val_accuracy: 0.9797
Epoch 11/50
1429/1429 [=====] - 7s 5ms/step - loss: 0.1448 - accuracy: 0.968
6 - val_loss: 0.2061 - val_accuracy: 0.9596
Epoch 12/50
1429/1429 [=====] - 7s 5ms/step - loss: 0.1496 - accuracy: 0.967
1 - val_loss: 0.1346 - val_accuracy: 0.9760
Epoch 13/50
1429/1429 [=====] - 7s 5ms/step - loss: 0.1034 - accuracy: 0.974
9 - val_loss: 0.1532 - val_accuracy: 0.9732
Epoch 14/50
1429/1429 [=====] - 7s 5ms/step - loss: 0.1316 - accuracy: 0.969
5 - val_loss: 0.1718 - val_accuracy: 0.9715
Epoch 15/50
1429/1429 [=====] - 7s 5ms/step - loss: 0.1541 - accuracy: 0.967
4 - val_loss: 0.1192 - val_accuracy: 0.9789
Epoch 16/50
1429/1429 [=====] - 7s 5ms/step - loss: 0.1088 - accuracy: 0.973
3 - val_loss: 0.1172 - val_accuracy: 0.9787
Epoch 17/50
1429/1429 [=====] - 7s 5ms/step - loss: 3.6909 - accuracy: 0.921
7 - val_loss: 0.4414 - val_accuracy: 0.9376
Epoch 18/50
1429/1429 [=====] - 7s 5ms/step - loss: 0.4344 - accuracy: 0.933
5 - val_loss: 0.3276 - val_accuracy: 0.9406
Epoch 19/50
1429/1429 [=====] - 7s 5ms/step - loss: 0.3686 - accuracy: 0.936
9 - val_loss: 0.3165 - val_accuracy: 0.9404
Epoch 20/50
1429/1429 [=====] - 7s 5ms/step - loss: 0.3356 - accuracy: 0.937
5 - val_loss: 0.2768 - val_accuracy: 0.9421
Epoch 21/50
1429/1429 [=====] - 7s 5ms/step - loss: 0.3583 - accuracy: 0.938
2 - val_loss: 0.4114 - val_accuracy: 0.9376
Epoch 22/50
1429/1429 [=====] - 7s 5ms/step - loss: 0.2693 - accuracy: 0.944
6 - val_loss: 0.2467 - val_accuracy: 0.9610
Epoch 23/50
1429/1429 [=====] - 7s 5ms/step - loss: 0.2832 - accuracy: 0.949
9 - val_loss: 0.3593 - val_accuracy: 0.9392
Epoch 24/50
1429/1429 [=====] - 7s 5ms/step - loss: 0.3753 - accuracy: 0.936
3 - val_loss: 0.2503 - val_accuracy: 0.9569
Epoch 25/50
1429/1429 [=====] - 7s 5ms/step - loss: 0.2299 - accuracy: 0.950
7 - val_loss: 0.2319 - val_accuracy: 0.9632
Epoch 26/50
1429/1429 [=====] - 7s 5ms/step - loss: 0.2041 - accuracy: 0.954
6 - val_loss: 0.2167 - val_accuracy: 0.9638
Epoch 27/50
1429/1429 [=====] - 7s 5ms/step - loss: 0.2001 - accuracy: 0.957
9 - val_loss: 0.2283 - val_accuracy: 0.9632
Epoch 28/50
1429/1429 [=====] - 7s 5ms/step - loss: 0.5929 - accuracy: 0.942
9 - val_loss: 0.3437 - val_accuracy: 0.9427
Epoch 29/50
1429/1429 [=====] - 7s 5ms/step - loss: 0.2959 - accuracy: 0.940
6 - val_loss: 0.2292 - val_accuracy: 0.9604
Epoch 30/50
1429/1429 [=====] - 7s 5ms/step - loss: 0.2613 - accuracy: 0.947
7 - val_loss: 0.2609 - val_accuracy: 0.9553
Epoch 31/50


```

Epoch 31/50
1429/1429 [=====] - 7s 5ms/step - loss: 0.2185 - accuracy: 0.952
9 - val_loss: 0.1931 - val_accuracy: 0.9654
Epoch 32/50
1429/1429 [=====] - 7s 5ms/step - loss: 0.1848 - accuracy: 0.958
0 - val_loss: 0.5031 - val_accuracy: 0.9419
Epoch 33/50
1429/1429 [=====] - 7s 5ms/step - loss: 0.2599 - accuracy: 0.947
8 - val_loss: 0.1827 - val_accuracy: 0.9673
Epoch 34/50
1429/1429 [=====] - 7s 5ms/step - loss: 0.1943 - accuracy: 0.957
2 - val_loss: 0.2247 - val_accuracy: 0.9634
Epoch 35/50
1429/1429 [=====] - 7s 5ms/step - loss: 0.1781 - accuracy: 0.962
5 - val_loss: 0.1714 - val_accuracy: 0.9717
Epoch 36/50
1429/1429 [=====] - 7s 5ms/step - loss: 0.1478 - accuracy: 0.967
8 - val_loss: 0.1683 - val_accuracy: 0.9736
Epoch 37/50
1429/1429 [=====] - 7s 5ms/step - loss: 0.1744 - accuracy: 0.963
8 - val_loss: 0.1670 - val_accuracy: 0.9671
Epoch 38/50
1429/1429 [=====] - 7s 5ms/step - loss: 0.1517 - accuracy: 0.965
2 - val_loss: 0.1655 - val_accuracy: 0.9744
Epoch 39/50
1429/1429 [=====] - 7s 5ms/step - loss: 0.1475 - accuracy: 0.967
8 - val_loss: 0.1666 - val_accuracy: 0.9665
Epoch 40/50
1429/1429 [=====] - 7s 5ms/step - loss: 0.1470 - accuracy: 0.967
0 - val_loss: 0.1607 - val_accuracy: 0.9726
Epoch 41/50
1429/1429 [=====] - 7s 5ms/step - loss: 3.9259 - accuracy: 0.926
6 - val_loss: 0.5860 - val_accuracy: 0.9376
Epoch 42/50
1429/1429 [=====] - 7s 5ms/step - loss: 0.4212 - accuracy: 0.932
5 - val_loss: 0.2822 - val_accuracy: 0.9567
Epoch 43/50
1429/1429 [=====] - 7s 5ms/step - loss: 0.2567 - accuracy: 0.946
7 - val_loss: 0.3919 - val_accuracy: 0.9427
Epoch 44/50
1429/1429 [=====] - 7s 5ms/step - loss: 0.2598 - accuracy: 0.946
3 - val_loss: 0.2666 - val_accuracy: 0.9646
Epoch 45/50
1429/1429 [=====] - 7s 5ms/step - loss: 0.2153 - accuracy: 0.954
1 - val_loss: 0.3079 - val_accuracy: 0.9490
Epoch 46/50
1429/1429 [=====] - 7s 5ms/step - loss: 0.2804 - accuracy: 0.945
6 - val_loss: 0.3669 - val_accuracy: 0.9535
Epoch 47/50
1429/1429 [=====] - 7s 5ms/step - loss: 0.2634 - accuracy: 0.950
0 - val_loss: 0.2404 - val_accuracy: 0.9657
Epoch 48/50
1429/1429 [=====] - 7s 5ms/step - loss: 0.2046 - accuracy: 0.957
4 - val_loss: 0.1993 - val_accuracy: 0.9622
Epoch 49/50
1429/1429 [=====] - 7s 5ms/step - loss: 0.1997 - accuracy: 0.957
7 - val_loss: 0.2294 - val_accuracy: 0.9642
Epoch 50/50
1429/1429 [=====] - 7s 5ms/step - loss: 0.2743 - accuracy: 0.952
5 - val_loss: 0.1802 - val_accuracy: 0.9719

```

Performance evaluation

In []:

```
print("Loss of the model is - " , model1.evaluate(X_test1,y_test1)[0])
```

```
print("Accuracy of the model is - " , model1.evaluate(X_test1,y_test1)[1]*100 , "%")
```

```
397/397 [=====] - 1s 3ms/step - loss: 0.1721 - accuracy: 0.9708  
Loss of the model is - 0.17208269238471985  
397/397 [=====] - 1s 3ms/step - loss: 0.1721 - accuracy: 0.9708  
Accuracy of the model is - 97.07897305488586 %
```

```
In [ ]:
```

```
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, cla  
ssification_report, confusion_matrix
```

```
In [ ]:
```

```
predict = model1.predict(X_test1)
```

```
In [ ]:
```

```
y_classes = [np.argmax(y, axis=None, out=None) for y in y_test1]
```

```
In [ ]:
```

```
yp_classes = [np.argmax(y, axis=None, out=None) for y in predict]
```

```
In [ ]:
```

```
accuracy_score(y_classes, yp_classes)
```

```
Out[ ]:
```

```
0.9707897015982994
```

```
In [ ]:
```

```
# micro-average will aggregate the contributions of all classes to compute the average me  
tric  
precision_score(y_classes, yp_classes, average='micro')
```

```
Out[ ]:
```

```
0.9707897015982994
```

```
In [ ]:
```

```
recall_score(y_classes, yp_classes, average='micro')
```

```
Out[ ]:
```

```
0.9707897015982994
```

```
In [ ]:
```

```
f1_score(y_classes, yp_classes, average='micro')
```

```
Out[ ]:
```

```
0.9707897015982994
```

```
In [ ]:
```

```
model1.save('hierogylyphs.h5')
```

```
In [ ]:
```

```
model1 = models.load_model('hierogylyphs.h5')
```

```
In [ ]:
```

```
print("Loss of the model is - " , model1.evaluate(rawImages2,labels2_one_hot_encoded)[0]
```

```
)  
print("Accuracy of the model is - " , model1.evaluate(rawImages2,labels2_one_hot_encoded  
) [1]*100 , "%")
```

```
62/62 [=====] - 2s 31ms/step - loss: 6175.4155 - accuracy: 0.006  
1  
Loss of the model is - 6175.41552734375  
62/62 [=====] - 2s 32ms/step - loss: 6175.4155 - accuracy: 0.006  
1  
Accuracy of the model is - 0.6100661121308804 %
```


Conclusion and Summary:

1. The dataset was loaded from google drive.
2. Feature extraction was done using *Inception Neural Network*.
3. SVM, XGBoost and MLV Classifiers were tarined and tested.

- SVM = 64% (Approx)
- XGBoost = 68% (Approx)
- MLV Classifier = 64% (Approx)

4. Some Examples were tested in the XGBoost Classifier.
5. Data was augmented for the neural network and manually preprocessed data was loaded.
6. A Custom CNN was defined and trained.

- Training accuracy = 95%
- Testing accuracy = 97%



As we see in the images, the two dataset have completely different image quality, the first image is a part of the image from the Piramid of Giza, where as the second image is a hand written image. So this explains why our model fails on the second dataset.